
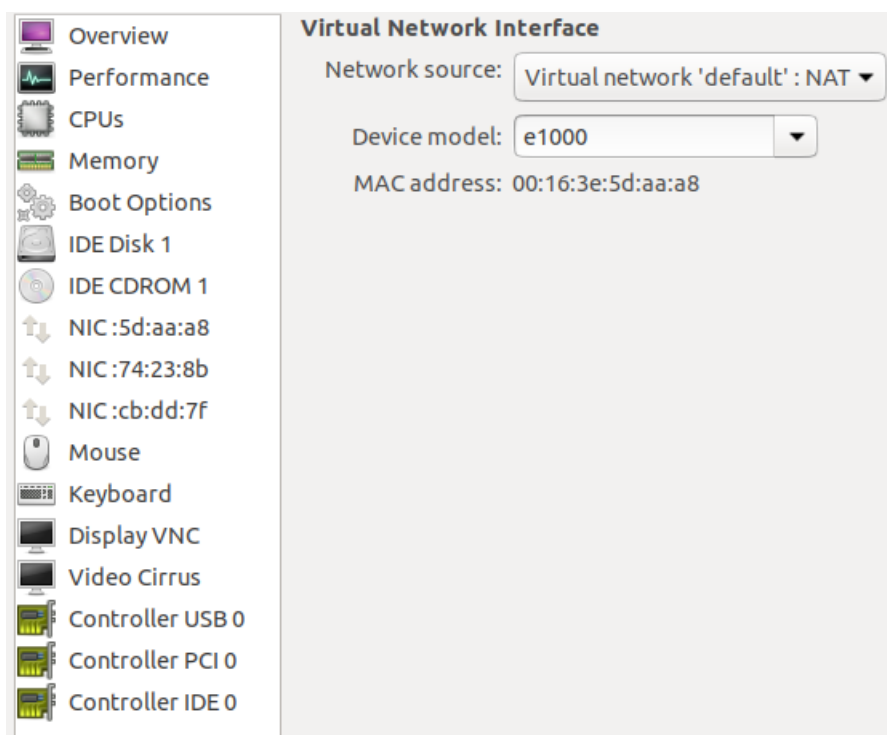



课程大作业 5

叶增渝 519030910168

1. 本实验需要两个 qemu 虚拟机分别进行 dpdk 的 l2fwd 执行与 pktgen-dpdk 的执行，为了防止环境配置冲突，所以我们选择重新创建两个 qemu 虚拟机。由于 dpdk 需要 Intel 的多网卡，但 QEMU 使用 xml 文件缺省为不支持 Intel 的单网卡，且由于 pktgen-dpdk 在运行时需要较多的 hugepage 内存，所以我们增大执行 pktgen-dpdk 的虚拟机所需的空间为 2GB，且为虚拟机配置多网卡，此处展示的是采用 virt-manager 的粗配置（网卡均采用 e1000）。此外我们还需要更改 xml 文件配置，详细内容在最后附带的源代码中。



两虚拟机如图所示（通过 virt-manager 打开）



对于虚拟机 VM1、VM2,我们分别运行 l2fwd 与 pktgen-dpdk, pktgen-dpdk 是基于 dpdk 的一个生成包的工具, 所以 VM2 需要在安装 dpdk 的基础上再安装 pktgen-dpdk, 对于 VM1, 我们仅需要照搬 VM2 安装 dpdk 的工作即可。(所需的所有文件如下所示)



2.首先需要安装 dpdk, 我们选择 19.11.10 这个 stable 版本将其解压到 desktop 下并改名为 dpdk 方便引用, 由于一些操作需要较高的权限, 所以我们选择全程使用 root 进行操作。我们首先需要安装的一些必要工具。

```
apt-get install python3
```

```
apt-get install numactl
```

```
apt-get install libnuma-dev
```

```
apt-get install net-tools
```

tips: python3 是为了让后续执行 setup 程序时部分结果能正常显示, net-tools 是为了了解当前虚拟机的 ip 与 mac 地址且对网卡进行一定的操作, 剩下两个则是安装需要的文件。

```
spoilvoid@spoilvoid:~/Desktop$ ifconfig
ens3: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.122.232  netmask 255.255.255.0  broadcast 192.168.122.255
    inet6 fe80::e35b:df95:59f8:52f5  prefixlen 64  scopeid 0x20<link>
    ether 00:16:3e:5d:aa:a9  txqueuelen 1000  (Ethernet)
    RX packets 822  bytes 897376 (897.3 KB)
    RX errors 64  dropped 0  overruns 0  frame 64
    TX packets 453  bytes 43289 (43.2 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

ens5: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.122.118  netmask 255.255.255.0  broadcast 192.168.122.255
    inet6 fe80::89f9:40a8:8396:6d3b  prefixlen 64  scopeid 0x20<link>
    ether 52:54:00:e5:20:55  txqueuelen 1000  (Ethernet)
    RX packets 82  bytes 10604 (10.6 KB)
    RX errors 63  dropped 0  overruns 0  frame 63
    TX packets 54  bytes 6185 (6.1 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

ens6: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
    inet 192.168.122.81  netmask 255.255.255.0  broadcast 192.168.122.255
    inet6 fe80::5d6c:41a6:ce51:d1c4  prefixlen 64  scopeid 0x20<link>
    ether 52:54:00:a3:60:31  txqueuelen 1000  (Ethernet)
    RX packets 82  bytes 10610 (10.6 KB)
    RX errors 62  dropped 0  overruns 0  frame 62
    TX packets 53  bytes 6101 (6.1 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING>  mtu 65536
    inet 127.0.0.1  netmask 255.0.0.0
    inet6 ::1  prefixlen 128  scopeid 0x10<host>
    loop txqueuelen 1000  (Local Loopback)
    RX packets 180  bytes 15467 (15.4 KB)
    RX errors 0  dropped 0  overruns 0  frame 0
    TX packets 180  bytes 15467 (15.4 KB)
    TX errors 0  dropped 0  overruns 0  carrier 0  collisions 0
```

此时虚拟机中的网卡上图所示, 为了将多出的网卡配置为 dpdk 兼容的网卡, 我们首先需要将其关闭

```
root@spoilvoid:/home/spoilvoid/Desktop# ifconfig ens5 down
root@spoilvoid:/home/spoilvoid/Desktop# ifconfig ens6 down
```

然后我们进行 dpdk 的编译安装，首先我们需要导入一些环境变量

```
root@spoilvoid:/home/spoilvoid/Desktop# export RTE_SDK=/home/spoilvoid/Desktop/dpdk
root@spoilvoid:/home/spoilvoid/Desktop# export RTE_TARGET=x86_64-native-linuxapp-gcc
```

然后调用 usertools 中的 dpdk-setup.sh 进行调配

```
root@spoilvoid:/home/spoilvoid/Desktop# cd /home/spoilvoid/Desktop/dpdk/usertools
root@spoilvoid:/home/spoilvoid/Desktop/dpdk/usertools# ./dpdk-setup.sh
-----
RTE_SDK exported as /home/spoilvoid/Desktop/dpdk
-----
Step 1: Select the DPDK environment to build
```

我们选择与上面导入的环境变量相应的 x86_64-native-linuxapp-gcc 选项进行编译

```
Option: 41

Configuration done using x86_64-native-linuxapp-gcc
== Build lib
== Build lib/librte_kvargs
== Build lib/librte_eal
== Build lib/librte_eal/common
== Build lib/librte_eal/linux
== Build lib/librte_eal/linux/eal
== Build lib/librte_pci
== Build lib/librte_ring
== Build lib/librte_stack
== Build lib/librte_mempool
== Build lib/librte_mbuf
== Build lib/librte_timer
== Build lib/librte_cfgfile
== Build lib/librte_net
== Build lib/librte_cmdline
== Build lib/librte_meter
== Build lib/librte_ethdev
== Build lib/librte_bbdev
== Build lib/librte_cryptodev
== Build lib/librte_security
== Build lib/librte_compressdev
== Build lib/librte_hash
== Build lib/librte_eventdev
== Build lib/librte_rawdev
== Build lib/librte_vhost
== Build lib/librte_efd
== Build lib/librte_rib
== Build lib/librte_fib
== Build lib/librte_lpm
== Build lib/librte_acl
== Build lib/librte_member
== Build lib/librte_ip_frag
== Build lib/librte_gro
== Build lib/librte_jobstats
```

在完成编译安装后导入 igb-uio 模块

```
Option: 48

Unloading any existing DPDK UIO module
Loading uio module
Loading DPDK UIO module
```

然后我们根据网卡的 PCI 地址配置两张之前被关闭的网卡（因为此操作在网卡 active 状态下

不可行)

```
Option: 54

Network devices using kernel driver
=====
0000:00:03.0 '82540EM Gigabit Ethernet Controller 100e' if=ens3 drv=e1000 unused=igb_uio,vfio-pci *Active*
0000:00:05.0 '82540EM Gigabit Ethernet Controller 100e' if=ens5 drv=e1000 unused=igb_uio,vfio-pci
0000:00:06.0 '82540EM Gigabit Ethernet Controller 100e' if=ens6 drv=e1000 unused=igb_uio,vfio-pci

No 'Baseband' devices detected
=====

No 'Crypto' devices detected
=====

No 'Eventdev' devices detected
=====

No 'Mempool' devices detected
=====

No 'Compress' devices detected
=====

No 'Misc (rawdev)' devices detected
=====

Enter PCI address of device to bind to IGB UIO driver: 00:05.0
OK
```

配置后的结果如下

```
Network devices using DPDK-compatible driver
=====
0000:00:05.0 '82540EM Gigabit Ethernet Controller 100e' drv=igb_uio unused=e1000,vfio-pci
0000:00:06.0 '82540EM Gigabit Ethernet Controller 100e' drv=igb_uio unused=e1000,vfio-pci

Network devices using kernel driver
=====
0000:00:03.0 '82540EM Gigabit Ethernet Controller 100e' if=ens3 drv=e1000 unused=igb_uio,vfio-pci *Active*

No 'Baseband' devices detected
=====

No 'Crypto' devices detected
=====

No 'Eventdev' devices detected
=====

No 'Mempool' devices detected
=====

No 'Compress' devices detected
=====

No 'Misc (rawdev)' devices detected
=====
```

最后我们需要配置 hugepage 内存

```
Option: 51

Removing currently reserved hugepages
Unmounting /mnt/huge and removing directory

Input the number of 2048kB hugepages
Example: to have 128MB of hugepages available in a 2MB huge page system,
enter '64' to reserve 64 * 2MB pages
Number of pages: 1024
Reserving hugepages
Creating /mnt/huge and mounting as hugetlbfs
```

配置后通过抓取巨页信息可以得到

```

root@spoilvoid:/home/spoilvoid/Desktop/dpdk/examples/l2fwd# grep -i huge /proc/meminfo
AnonHugePages:          0 kB
ShmemHugePages:         0 kB
FileHugePages:          0 kB
HugePages_Total:       513
HugePages_Free:        512
HugePages_Rsvd:         0
HugePages_Surp:         0
Hugepagesize:          2048 kB
Hugetlb:               1050624 kB

```

到此处, dpdk 的安装就完成了, 然后我们进入 examples 文件夹对我们接下来要执行的 l2fwd 进行 make 操作, 获取可执行文件

```

root@spoilvoid:/home/spoilvoid/Desktop/dpdk# cd examples
root@spoilvoid:/home/spoilvoid/Desktop/dpdk/examples# cd l2fwd
root@spoilvoid:/home/spoilvoid/Desktop/dpdk/examples/l2fwd# make
CC main.o
LD l2fwd
INSTALL-APP l2fwd
INSTALL-MAP l2fwd.map

```

3.接下来我们需要安装 pktgen-dpdk, 由前面图片可知我们选择的 19.12 版本的 pktgen-dpdk, 同样将其解压到 Desktop 文件夹下, 并命名为 pktgen-dpdk, 但是这里我们并没有办法直接进行 make 安装, 因为缺少一些工具, 我们需要先下载必要工具:

apt-get install libpcap-dev

apt-get install libreadline-dev

然后我们需要手动安装 lua, 解压 lua5.3.5 文件, 并进入对应的文件夹内, 使用 Makefile 指定的 make linux install 命令进行安装

```

root@spoilvoid:/home/spoilvoid/Downloads/lua-5.3.5# su
Password:
root@spoilvoid:/home/spoilvoid/Downloads/lua-5.3.5# make linux install
cd src && make linux
make[1]: Entering directory '/home/spoilvoid/Downloads/lua-5.3.5/src'
make all SYSCFLAGS="-DLUA_USE_LINUX" SYSLIBS="-lreadline"
make[2]: Entering directory '/home/spoilvoid/Downloads/lua-5.3.5/src'
make[2]: Nothing to be done for 'all'.
make[2]: Leaving directory '/home/spoilvoid/Downloads/lua-5.3.5/src'
make[1]: Leaving directory '/home/spoilvoid/Downloads/lua-5.3.5/src'
cd src && mkdir -p /usr/local/bin /usr/local/include /usr/local/lib /usr/local/man/man1 /usr/local/share/lua/5.3 /usr/local/lib/lua/5.3
cd src && install -p -m 0755 lua luac /usr/local/bin
cd src && install -p -m 0644 lua.h luaconf.h luaLlib.h lauxlib.h lua.hpp /usr/local/include
cd src && install -p -m 0644 liblua.a /usr/local/lib
cd doc && install -p -m 0644 lua.1 luac.1 /usr/local/man/man1
root@spoilvoid:/home/spoilvoid/Downloads/lua-5.3.5#

```

在此之后我们进入 pktgen-dpdk 文件夹下正式开始 pktgen-dpdk 的安装 (这里需要保持环境变量与上述相同), 进入相应文件夹下执行 make 即可

```

root@spoilvoid:/home/spoilvoid/Desktop/dpdk/usertools# cd /home/spoilvoid/Desktop/pktgen-dpdk
root@spoilvoid:/home/spoilvoid/Desktop/pktgen-dpdk# make
== lib
== common
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/common/ABI_VERSION: No such file or directory
== plugin
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/plugin/ABI_VERSION: No such file or directory
== utils
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/utils/ABI_VERSION: No such file or directory
== vec
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/vec/ABI_VERSION: No such file or directory
== lua
"Lua pkg-config file was not found"
"Get lua-5.3 from lua.org and build it on your system"
"Also install lua 5.3 into /usr/local/include and /usr/local/lib"
"Make sure the library in /usr/local/lib is called liblua.a"
"Lua pkg-config file was not found"
"Get lua-5.3 from lua.org and build it on your system"
"Also install lua 5.3 into /usr/local/include and /usr/local/lib"
"Make sure the library in /usr/local/lib is called liblua.a"
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/lua/ABI_VERSION: No such file or directory
== cli
cat: /home/spoilvoid/Desktop/pktgen-dpdk/lib/cli/ABI_VERSION: No such file or directory
== app
"Lua pkg-config was not found"
"Get lua-5.3 from lua.org and build it on your system"
"Also install lua 5.3 into /usr/local/include and /usr/local/lib"
"Make sure the library in /usr/local/lib is called liblua.a"
"Lua pkg-config was not found"
"Get lua-5.3 from lua.org and build it on your system"
"Also install lua 5.3 into /usr/local/include and /usr/local/lib"
"Make sure the library in /usr/local/lib is called liblua.a"
LD pktgen
INSTALL-APP pktgen
INSTALL-MAP pktgen.map

```

4.此时 VM2 所需的环境与工具已经配置完毕, VM1 只需要按照上述的 dpdk 配置步骤编译安装即可, 接下来为了使得 pktgen-dpdk 的 VM2 能够向 VM1 发包, 我们需要得知两 qemu 虚拟机的 ip 地址与 VM1 的两个端口的 mac 地址 (一个可通过 virsh 指令得到, 一个可以在执行 l2fwd 时获得)

```

virsh # list

```

Id	Name	State
1	test2	running
2	test1	running

```

virsh # domifaddr 1

```

Name	MAC address	Protocol	Address
vnet0	00:16:3e:5d:aa:a9	ipv4	192.168.122.232/24

```

virsh # domifaddr 2

```

Name	MAC address	Protocol	Address
vnet3	00:16:3e:5d:aa:a8	ipv4	192.168.122.231/24


```

root@spoilvoid:/home/spoilvoid/Desktop/dpdk/examples/l2fwd/build# ./l2fwd -c 0x3 -n 2 -- -p 3 -q 1
EAL: Detected 2 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:00:03.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:05.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:06.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
MAC updating enabled
Lcore 0: RX port 0
Lcore 1: RX port 1
Initializing port 0... done:
Port 0, MAC address: 52:54:00:74:23:8B

Initializing port 1... done:
Port 1, MAC address: 52:54:00:CB:DD:7F

Checking link status.....done
Port0 Link Up. Speed 1000 Mbps - full-duplex
Port1 Link Up. Speed 1000 Mbps - full-duplex
L2FWD: entering main loop on lcore 1
L2FWD: -- lcoreid=1 portid=1
L2FWD: entering main loop on lcore 0
L2FWD: -- lcoreid=0 portid=0

```

5.接下来我们开始正式进行 VM2 生成包，并发送到 VM1 的 0 号端口上，VM1 的 1 号端口则会转发从 0 号端口进入的包，而这个包会重新被 VM2 接收，从而产生一个接受速率与发送速率，由此进行性能比较

情况 1: 仅在 VM1 上执行 l2fwd

命令: ./l2fwd -c 0x3 -n 2 -- -p 3 -q 1

```

root@spoilvoid:/home/spoilvoid/Desktop/dpdk/examples/l2fwd/build# ./l2fwd -c 0x3 -n 2 -- -p 3 -q 1
EAL: Detected 2 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:00:03.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:05.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:06.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
MAC updating enabled
Lcore 0: RX port 0
Lcore 1: RX port 1
Initializing port 0... done:
Port 0, MAC address: 52:54:00:74:23:8B

Initializing port 1... done:
Port 1, MAC address: 52:54:00:CB:DD:7F

Checking link status.....done
Port0 Link Up. Speed 1000 Mbps - full-duplex
Port1 Link Up. Speed 1000 Mbps - full-duplex
L2FWD: entering main loop on lcore 1
L2FWD: -- lcoreid=1 portid=1
L2FWD: entering main loop on lcore 0
L2FWD: -- lcoreid=0 portid=0

```

此时由于没有外界发包，仅仅是端口混杂模式互相转发，所以两个端口的 sent 与 received 量相同

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                0
Packets received:            0
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                0
Packets received:            0
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          0
Total packets received:      0
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                552253
Packets received:            582414
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                582414
Packets received:            552253
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          1134667
Total packets received:      1134667
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                1602127
Packets received:            1697472
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                1697472
Packets received:            1602127
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          3299599
Total packets received:      3299599
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                2097984
Packets received:            2237907
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                2237907
Packets received:            2097984
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          4335891
Total packets received:      4335891
Total packets dropped:        0
=====

```



```

Port statistics =====
Statistics for port 0 -----
Packets sent:                3057249
Packets received:            3307668
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                3307668
Packets received:            3057281
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          6364949
Total packets received:      6364981
Total packets dropped:        0
=====

```

情况 2：仅在 VM2 上执行 pktgen

命令：./pktgen -l 0-1 -n 3 -- -P -m "[1].0"

```

root@spoilvoid: /home/spoilvoid/Desktop/pktgen-dpdk/app/x86_64-native-linuxapp-gcc/app# ./pktgen -l 0-1 -
n 3 -- -P -m "[1].0"

Copyright (c) <2010-2019>, Intel Corporation. All rights reserved. Powered by DPDK
EAL: Detected 2 lcore(s)
EAL: Detected 1 NUMA nodes
EAL: Multi-process socket /var/run/dpdk/rte/mp_socket
EAL: Selected IOVA mode 'PA'
EAL: Probing VFIO support...
EAL: VFIO support initialized
EAL: PCI device 0000:00:03.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:05.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
EAL: PCI device 0000:00:06.0 on NUMA socket -1
EAL:   probe driver: 8086:100e net_e1000_em
Lua 5.3.5 Copyright (C) 1994-2018 Lua.org, PUC-Rio

*** Copyright (c) <2010-2019>, Intel Corporation. All rights reserved.
*** Pktgen created by: Keith Wiles -- >>> Powered by DPDK <<<

Port: Name          IfIndex Alias          NUMA  PCI
   0: net_e1000_em    0              0      8086:100e/00:05.0
   1: net_e1000_em    0              0      8086:100e/00:06.0

Initialize Port 0 -- TxQ 1, RxQ 1
Src MAC 52:54:00:e5:20:55
<Promiscuous mode Enabled>

RX/TX processing lcore:   1 rx:  1 tx:  1

```

初始状态什么都不做时，观察下图不难发现发送速率与接收速率均为 0（这里由于放大缩小产生了错位）

```

Port: Name          I P-----Single          :0MA PCI
  0: net_e1000_em    <UP-1000-FD>0      8086:100e/00:05.0
  1: net_e1000_em    153/0                153/0
                        0/0                0/0
Initialize Port 0 -- 0/0                0/0
Src MAC 52:54:00:e5:2 0
<Promiscuous mode En 10
                        174
                        100
RX/TX processing lc  16
                        12
- Ports 0-1 of 2 <M 0c) <2010-2019>, Intel Corporation
  Flags:Port         : 0
-link State          : 0/0      ---Total Rate---
Pkts/s Max/Rx        : 3/0      0/0
      Max/Tx         : 0/0      0/0
Mbits/s Rx/Tx        : 301      153/1
Broadcast            : 0        0/0
Multicast             : 0        0/0
Sizes 64              : 0
      65-127         : 0
      128-255        : 147
      256-511        : 11
      512-1023       : 2
      1024-1518      : 0
Runts/Jumbos         : 0
ARP/ICMP Pkts        : 0
Errors Rx/Tx         : 0/0
Total Rx Pkts        : 0/0
      Tx Pkts         : 0/0
      Rx MBs          : 160
      Tx MBs          : 0
                        : 0
Pattern Type         : 0

```

当我们执行 start 0 命令时，由于没有正确地设置目的 ip、mac，所以仅有发送速率没有接受速率

```

/nPorts 0-1 of 20 -<Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0              0/0
Link State      :      <UP-1000-FD>      ---Total Rate---
Pkts/s Max/Rx   :      153/1              153/1
      Max/Tx     :      520000/514944      520000/514944
Mbits/s Rx/Tx   :      0/346              0/346
Broadcast       :      0
Multicast       :      11c) <2010-2019>, Intel Corporation
Sizes 64        :      260
      65-127     :      129      ---Total Rate---
      128-255    :      19              0/0
      256-511    :      12c) <2010-2019>, Intel Corporation
      512-1023   :      0c) <2010-2019>, Intel Corporation
      1024-1518  :      0c) <2010-2019>, Intel Corporation
Runts/Jumbos    :      0/0      ---Total Rate---
ARP/ICMP Pkts   :      4/0      ---Total Rate---
Errors Rx/Tx    :      0/0              153/1
Total Rx Pkts   :      420              0/0
      Tx Pkts    :      2889536          153/0
      Rx MBs     :      0              0/0
      Tx MBs     :      1942             0/0
Sizes 64        :      0
Pattern Type     :      abcd...
Tx Count/% Rate  :      Forever /100%
Pkt Size/Tx Burst :      64 / 64
TTL/Port Src/Dest :      4/ 1234/ 5678
Pkt Type:VLAN ID :      IPv4 / TCP:0001
802.1p CoS/DSCP/IPP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0
IP Destination   :      192.168.122.231
      Source     :      192.168.122.232/24by DPDK (pid:38134) -----
MAC Destination  :      52:54:00:74:23:8b
      Source     :      52:54:00:e5:20:55erface without timers
PCI Vendor/Addr  :      8086:100e/00:05.0
      Tx MBs     :      0
-- Pktgen 19.12.0 (DPDK 19.11.10) Powered by DPDK (pid:38134) -----
Pktgen:/> start 0
Pktgen:/>

```

情况 3:

在 VM1 上执行 l2fwd

命令: ./l2fwd -c 0x3 -n 2 -- -p 3 -q 1

在 VM2 上执行 pktgen

命令: ./pktgen -l 0-1 -n 3 -- -P -m "[1].0"

并进行 dst ip、mac 与 src ip 的设置 (src mac 已经设置好了, 由于命令无法顺序展示, 所以以这种方式展示) 其中 start 0 为开始发包指令, 对应停止发包指令为 stop 0

```

set 0 dst ip 192.168.122.231
set 0 src ip 192.168.122.232/24
set 0 dst mac 52:54:00:74:23:8B
start 0

```

然后我们首先观察 pktgen, 可以看到发送与接收速率分别为 18~19MB/s 与 8~9MB/s, 且发送指向了正确的地址 (因为我们正确地所接收到了转发回来的包)

```

\ Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0          126436/96705
Link State      : <UP-1000-FD>        ---Total Rate---
Pkts/s Max/Rx   : 126436/99783        126436/99783
Max/Tx          : 556352/224512        556352/224512
Mbits/s Rx/Tx   : 70/150              70/150
Broadcast       : 0                   0/0
Multicast       : 11c) <2010-2019>, Intel Corporation
Sizes 64        : 8985350c) <2010-2019>, Intel Corporation
65-127          : 132530c) <2010-2019>, Intel Corporation
128-255         : 19                  ---Total Rate---
256-511         : 12                  ---Total Rate---
512-1023        : 0                   153/1
1024-1518       : 0                   0/0
Runts/Jumbos    : 0/0                153/0
ARP/ICMP Pkts   : 4/0c) <2010-2019>, Intel Corporation
Errors Rx/Tx    : 0/0                0/0
Total Rx Pkts   : 9040696             126436/1
Tx Pkts         : 25298304            ---Total Rate---
Rx MBs          : 6403                126436/1
Tx MBs          : 17016               520000/453056
Broadcast       : 0/304              0/304
Pattern Type    : abcd...
Tx Count/% Rate : Forever /100%
Pkt Size/Tx Burst : 64 / 64
TTL/Port Src/Dest : 4/ 1234/ 5678
Pkt Type:VLAN ID : IPv4 / TCP:0001by DPDK (pid:38134) ----- 512-1023
802.1p CoS/DSCP/IPP : 0/ 0/ 0
VxLAN Flg/Grp/vid : 0000/ 0/ 0erface without timers
IP Destination   : 192.168.122.231
Source           : 192.168.122.232/24
MAC Destination  : 52:54:00:74:23:8bby DPDK (pid:38134) -----Total Rx Pkts
Source           : 52:54:00:e5:20:55
PCI Vendor/Addr  : 8086:100e/00:05.0
Rx MBs           : 18772992
-- Pktgen 19.12.0 (DPDK 19.11.10) Powered by DPDK (pid:38134) -----

\ Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0          126436/96705
Link State      : <UP-1000-FD>        ---Total Rate---
Pkts/s Max/Rx   : 126436/109426       126436/109426
Max/Tx          : 556352/229056       556352/229056
Mbits/s Rx/Tx   : 77/153              77/153
Broadcast       : 0                   0/0
Multicast       : 11c) <2010-2019>, Intel Corporation
Sizes 64        : 10425586c) <2010-2019>, Intel Corporation
65-127          : 132530c) <2010-2019>, Intel Corporation
128-255         : 19                  ---Total Rate---
256-511         : 12                  ---Total Rate---
512-1023        : 0                   153/1
1024-1518       : 0                   0/0
Runts/Jumbos    : 0/0                153/0
ARP/ICMP Pkts   : 4/0c) <2010-2019>, Intel Corporation
Errors Rx/Tx    : 0/0                0/0
Total Rx Pkts   : 10522602            126436/1
Tx Pkts         : 28400832            ---Total Rate---
Rx MBs          : 7447                126436/1
Tx MBs          : 19102               520000/453056
Broadcast       : 0/304              0/304
Pattern Type    : abcd...
Tx Count/% Rate : Forever /100%
Pkt Size/Tx Burst : 64 / 64
TTL/Port Src/Dest : 4/ 1234/ 5678
Pkt Type:VLAN ID : IPv4 / TCP:0001by DPDK (pid:38134) ----- 512-1023
802.1p CoS/DSCP/IPP : 0/ 0/ 0
VxLAN Flg/Grp/vid : 0000/ 0/ 0erface without timers
IP Destination   : 192.168.122.231
Source           : 192.168.122.232/24
MAC Destination  : 52:54:00:74:23:8bby DPDK (pid:38134) -----Total Rx Pkts
Source           : 52:54:00:e5:20:55
PCI Vendor/Addr  : 8086:100e/00:05.0
Rx MBs           : 18772992
-- Pktgen 19.12.0 (DPDK 19.11.10) Powered by DPDK (pid:38134) -----

```

当我们键入命令 stop 0 时，继续观察，可以看到虽然没有继续发包了，但是由于先前接受速率小于发送速率，VM1 依旧在转发这些数据包给我们，所以接受速率暂时还不为 0

```

/ Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0      126436/96705
Link State      :      <UP-1000-FD>      ---Total Rate---
Pkts/s Max/Rx   :      127900/119562      127900/119562
Max/Tx          :      556352/0      556352/0
Mbits/s Rx/Tx   :      84/0      84/0
Broadcast       :      0      0/0
Multicast       :      11c) <2010-2019>, Intel Corporation
Sizes 64        :      15715665c) <2010-2019>, Intel Corporation
65-127         :      132530c) <2010-2019>, Intel Corporation
128-255        :      19      ---Total Rate---
256-511        :      12      ---Total Rate---
512-1023       :      0      153/1
1024-1518      :      0      0/0
Runts/Jumbos    :      0/0      153/0
ARP/ICMP Pkts   :      4/0c) <2010-2019>, Intel Corporation
Errors Rx/Tx    :      0/0      0/0
Total Rx Pkts   :      15708228      126436/1
Tx Pkts         :      36170496      ---Total Rate---
Rx MBs          :      11103      126436/1
Tx MBs          :      24333      520000/453056
Broadcast       :      0/304      0/304
Pattern Type    :      abcd...
Tx Count/% Rate :      Forever /100%
Pkt Size/Tx Burst :      64 / 64
TTL/Port Src/Dest :      4/ 1234/ 5678
Pkt Type:VLAN ID :      IPv4 / TCP:0001by DPDK (pid:38134) ----- 512-1023
802.1p CoS/DSCP/IPP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0erface without timers
IP Destination   :      192.168.122.231
Source           :      192.168.122.232/24
MAC Destination  :      52:54:00:74:23:8bby DPDK (pid:38134) -----Total Rx Pkts
Source           :      52:54:00:e5:20:55
PCI Vendor/Addr  :      8086:100e/00:05.0
Rx MBs           :      18772992
-- Pktgen 19.12.0 (DPDK 19.11.10) Powered by DPDK (pid:38134) -----
Pktgen:/> stop 0

```

```

| Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0      126436/96705
Link State      :      <UP-1000-FD>      ---Total Rate---
Pkts/s Max/Rx   :      129137/117241      129137/117241
Max/Tx          :      556352/0      556352/0
Mbits/s Rx/Tx   :      82/0      82/0
Broadcast       :      0      0/0
Multicast       :      11c) <2010-2019>, Intel Corporation
Sizes 64        :      17771392c) <2010-2019>, Intel Corporation
65-127         :      132530c) <2010-2019>, Intel Corporation
128-255        :      19      ---Total Rate---
256-511        :      12      ---Total Rate---
512-1023       :      0      153/1
1024-1518      :      0      0/0
Runts/Jumbos    :      0/0      153/0
ARP/ICMP Pkts   :      4/0c) <2010-2019>, Intel Corporation
Errors Rx/Tx    :      0/0      0/0
Total Rx Pkts   :      17838440      126436/1
Tx Pkts         :      36170496      ---Total Rate---
Rx MBs          :      12603      126436/1
Tx MBs          :      24333      520000/453056
Broadcast       :      0/304      0/304
Pattern Type    :      abcd...
Tx Count/% Rate :      Forever /100%
Pkt Size/Tx Burst :      64 / 64
TTL/Port Src/Dest :      4/ 1234/ 5678
Pkt Type:VLAN ID :      IPv4 / TCP:0001by DPDK (pid:38134) ----- 512-1023
802.1p CoS/DSCP/IPP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0erface without timers
IP Destination   :      192.168.122.231
Source           :      192.168.122.232/24
MAC Destination  :      52:54:00:74:23:8bby DPDK (pid:38134) -----Total Rx Pkts
Source           :      52:54:00:e5:20:55
PCI Vendor/Addr  :      8086:100e/00:05.0
Rx MBs           :      18772992
-- Pktgen 19.12.0 (DPDK 19.11.10) Powered by DPDK (pid:38134) -----
Pktgen:/> stop 0

```

但在 VM1 处理完所有发来的数据包后，接收速率也降为零。

```
\ Ports 0-1 of 2 <Main Page> Copyright (c) <2010-2019>, Intel Corporation
Flags:Port      : P-----Single      :0      126436/96705
Link State      :      <UP-1000-FD>      ---Total Rate---
Pkts/s Max/Rx   :      129137/1      129137/1
      Max/Tx     :      556352/0      556352/0
Mbits/s Rx/Tx   :      0/0      0/0
Broadcast       :      0      0/0
Multicast       :      11c) <2010-2019>, Intel Corporation
Sizes 64        :      20083520c) <2010-2019>, Intel Corporation
      65-127     :      132530c) <2010-2019>, Intel Corporation
      128-255    :      19      ---Total Rate---
      256-511    :      12      ---Total Rate---
      512-1023   :      0      153/1
      1024-1518  :      0      0/0
Runts/Jumbos    :      0/0      153/0
ARP/ICMP Pkts   :      4/0c) <2010-2019>, Intel Corporation
Errors Rx/Tx    :      0/0      0/0
Total Rx Pkts   :      20152801      126436/1
      Tx Pkts    :      36170496      ---Total Rate---
      Rx MBs     :      14234      126436/1
      Tx MBs     :      24333      520000/453056
Broadcast       :      0/304      0/304
Pattern Type    :      abcd...
Tx Count/% Rate :      Forever /100%
Pkt Size/Tx Burst :      64 / 64
TTL/Port Src/Dest :      4/ 1234/ 5678
Pkt Type:VLAN ID :      IPv4 / TCP:0001by DPKD (pid:38134) ----- 512-1023
802.1p CoS/DSCP/IP :      0/ 0/ 0
VxLAN Flg/Grp/vid :      0000/ 0/ 0erface without timers
IP Destination  :      192.168.122.231
      Source     :      192.168.122.232/24
MAC Destination :      52:54:00:74:23:8bby DPKD (pid:38134) -----Total Rx Pkts
      Source     :      52:54:00:e5:20:55
PCI Vendor/Addr :      8086:100e/00:05.0
      Rx MBs     :      18772992
-- Pktgen 19.12.0 (DPKD 19.11.10) Powered by DPKD (pid:38134) -----
Pktgen:/> stop 0
```

我们再观察观察 VM1 上的情况，可以看到 0 号端口接受的包数量远远超过了发送数量，port1 发送数量超过了接受数量，这是由于我们向 VM1 的零号端口发包导致的，而我们可以看到 VM2 也有接受速率，所以整体的实验成功了（之所以 0 号端口还存在发送包，1 号端口还存在接受包是由于开启了 -p 混杂模式）

```
Port statistics =====
Statistics for port 0 -----
Packets sent:      0
Packets received:  0
Packets dropped:   0
Statistics for port 1 -----
Packets sent:      0
Packets received:  0
Packets dropped:   0
Aggregate statistics =====
Total packets sent:      0
Total packets received:  0
Total packets dropped:   0
=====
```



```

Port statistics =====
Statistics for port 0 -----
Packets sent:                283081
Packets received:            766637
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                766637
Packets received:            283104
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          1049718
Total packets received:      1049741
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                587189
Packets received:            1514218
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                1514218
Packets received:            587218
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          2101407
Total packets received:      2101436
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                906613
Packets received:            2128609
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                2128609
Packets received:            906645
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          3035222
Total packets received:      3035254
Total packets dropped:        0
=====

```

```

Port statistics =====
Statistics for port 0 -----
Packets sent:                1248455
Packets received:            2695295
Packets dropped:              0
Statistics for port 1 -----
Packets sent:                2695295
Packets received:            1248455
Packets dropped:              0
Aggregate statistics =====
Total packets sent:          3943750
Total packets received:      3943750
Total packets dropped:        0
=====

```

6.我们成功地测试了传输速率与接受速率，但是随之而来产生了一个问题：我们发现在传输

过程中没有产生丢包

猜测原因：这可能是由于如果在网卡配置过程中使用 virtio，那么运行 l2fwd 会存在自收发机制，将多余的超过容量的包抛弃（类似于网络中缓冲区队列被占满导致丢包），但是由于我们将网卡设置为 e1000 且将 0 号端口设置为接收，1 号端口用于发包（两者之间还有额外的发送接收是因为开启了混杂模式），那么这种情况下没有自收发机制，完全地进行转运使得丢包率变为 0

7.最终性能结果：在使用 e1000 双端口的情况下，发送与接收速率分别为 18~19MB/s 与 8~9MB/s，丢包率为 0

配置 xml 文件源文件 test1.xml:

```
<domain type='kvm' id='1'>
  <name>test1</name>
  <uuid>493ce2c4-e75d-4d9c-be91-e2a6b41ca7d6</uuid>
  <memory unit='KiB'>1048576</memory>
  <currentMemory unit='KiB'>1048576</currentMemory>
  <vcpu placement='static'>2</vcpu>
  <cpu match='exact'>
    <model>Haswell</model>
    <feature policy='force' name='x2apic'/>
    <feature policy='force' name='pdpe1gb'/>
    <feature policy='disable' name='hle'/>
    <feature policy='disable' name='smep'/>
    <feature policy='disable' name='rtm'/>
  </cpu>
  <resource>
    <partition>/machine</partition>
  </resource>
  <os>
    <type arch='x86_64' machine='pc-i440fx-2.11'>hvm</type>
    <boot dev='hd'/>
  </os>
  <features>
    <acpi/>
    <apic/>
    <pae/>
  </features>
  <clock offset='localtime'/>
  <on_poweroff>destroy</on_poweroff>
  <on_reboot>restart</on_reboot>
  <on_crash>destroy</on_crash>
  <devices>
    <emulator>/usr/bin/kvm</emulator>
    <disk type='file' device='disk'>
```

```

    <driver name='qemu' type='qcow2'/>
    <source file='/home/spoilvoid/Desktop/3D/test1.img'/>
    <backingStore/>
    <target dev='hda' bus='ide'/>
    <alias name='ide0-0-0'/>
    <address type='drive' controller='0' bus='0' target='0' unit='0'/>
</disk>
<disk type='file' device='cdrom'>
    <driver name='qemu' type='raw'/>
    <source file='/home/spoilvoid/Desktop/3D/ubuntu-20.04.3-desktop-amd64.iso'/>
    <backingStore/>
    <target dev='hdb' bus='ide'/>
    <readonly/>
    <alias name='ide0-0-1'/>
    <address type='drive' controller='0' bus='0' target='0' unit='1'/>
</disk>
<controller type='usb' index='0' model='piix3-uhci'>
    <alias name='usb'/>
    <address type='pci' domain='0x0000' bus='0x00' slot='0x01' function='0x2'/>
</controller>
<controller type='pci' index='0' model='pci-root'>
    <alias name='pci.0'/>
</controller>
<controller type='ide' index='0'>
    <alias name='ide'/>
    <address type='pci' domain='0x0000' bus='0x00' slot='0x01' function='0x1'/>
</controller>
<interface type='network'>
    <mac address='00:16:3e:5d:aa:a8'/>
    <source network='default' bridge='virbr0'/>
    <target dev='vnet0'/>
    <model type='e1000'/>
    <alias name='net0'/>
    <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0'/>
</interface>
<interface type='network'>
    <mac address='52:54:00:74:23:8b'/>
    <source network='default' bridge='virbr0'/>
    <target dev='vnet1'/>
    <model type='e1000'/>
    <alias name='net1'/>
    <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x0'/>
</interface>
<interface type='network'>

```

```

    <mac address='52:54:00:cb:dd:7f'/>
    <source network='default' bridge='virbr0'/>
    <target dev='vnet2'/>
    <model type='e1000'/>
    <alias name='net2'/>
    <address type='pci' domain='0x0000' bus='0x00' slot='0x06' function='0x0'/>
</interface>
<input type='mouse' bus='ps2'>
  <alias name='input0'/>
</input>
<input type='keyboard' bus='ps2'>
  <alias name='input1'/>
</input>
<graphics type='vnc' port='5900' autoport='yes' listen='0.0.0.0' keymap='en-us'>
  <listen type='address' address='0.0.0.0'/>
</graphics>
<video>
  <model type='cirrus' vram='16384' heads='1' primary='yes'/>
  <alias name='video0'/>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x02' function='0x0'/>
</video>
<memballoon model='virtio'>
  <alias name='balloon0'/>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x04' function='0x0'/>
</memballoon>
</devices>
<seclabel type='dynamic' model='apparmor' relabel='yes'>
  <label>libvirt-493ce2c4-e75d-4d9c-be91-e2a6b41ca7d6</label>
  <imagelabel>libvirt-493ce2c4-e75d-4d9c-be91-e2a6b41ca7d6</imagelabel>
</seclabel>
<seclabel type='dynamic' model='dac' relabel='yes'>
  <label>+64055:+127</label>
  <imagelabel>+64055:+127</imagelabel>
</seclabel>
</domain>

```

配置 xml 文件源文件 test2.xml:

```

<domain type='kvm' id='2'>
  <name>test2</name>
  <uuid>faf9ab4e-57d0-4363-8ed2-15b5326918a5</uuid>
  <memory unit='KiB'>1048576</memory>
  <currentMemory unit='KiB'>1048576</currentMemory>
  <vcpu placement='static'>2</vcpu>

```

```

<cpu match='exact'>
  <model>Haswell</model>
  <feature policy='force' name='x2apic'/>
  <feature policy='force' name='pdpe1gb'/>
  <feature policy='disable' name='hle'/>
  <feature policy='disable' name='smep'/>
  <feature policy='disable' name='rtm'/>
</cpu>
<resource>
  <partition>/machine</partition>
</resource>
<os>
  <type arch='x86_64' machine='pc-i440fx-2.11'>hvm</type>
  <boot dev='hd'/>
</os>
<features>
  <acpi/>
  <apic/>
  <pae/>
</features>
<clock offset='localtime'/>
<on_poweroff>destroy</on_poweroff>
<on_reboot>restart</on_reboot>
<on_crash>destroy</on_crash>
<devices>
  <emulator>/usr/bin/kvm</emulator>
  <disk type='file' device='disk'>
    <driver name='qemu' type='qcow2'/>
    <source file='/home/spoilvoid/Desktop/3D/test2.img'/>
    <backingStore/>
    <target dev='hda' bus='ide'/>
    <alias name='ide0-0-0'/>
    <address type='drive' controller='0' bus='0' target='0' unit='0'/>
  </disk>
  <disk type='file' device='cdrom'>
    <driver name='qemu' type='raw'/>
    <source file='/home/spoilvoid/Desktop/3D/ubuntu-20.04.3-desktop-amd64.iso'/>
    <backingStore/>
    <target dev='hdb' bus='ide'/>
    <readonly/>
    <alias name='ide0-0-1'/>
    <address type='drive' controller='0' bus='0' target='0' unit='1'/>
  </disk>
  <controller type='usb' index='0' model='piix3-uhci'>

```

```

    <alias name='usb' />
    <address type='pci' domain='0x0000' bus='0x00' slot='0x01' function='0x2' />
</controller>
<controller type='pci' index='0' model='pci-root'>
    <alias name='pci.0' />
</controller>
<controller type='ide' index='0'>
    <alias name='ide' />
    <address type='pci' domain='0x0000' bus='0x00' slot='0x01' function='0x1' />
</controller>
<interface type='network'>
    <mac address='00:16:3e:5d:aa:a9' />
    <source network='default' bridge='virbr0' />
    <target dev='vnet0' />
    <model type='e1000' />
    <alias name='net0' />
    <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0' />
</interface>
<interface type='network'>
    <mac address='52:54:00:e5:20:55' />
    <source network='default' bridge='virbr0' />
    <target dev='vnet1' />
    <model type='e1000' />
    <alias name='net1' />
    <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x0' />
</interface>
<interface type='network'>
    <mac address='52:54:00:a3:60:31' />
    <source network='default' bridge='virbr0' />
    <target dev='vnet2' />
    <model type='e1000' />
    <alias name='net2' />
    <address type='pci' domain='0x0000' bus='0x00' slot='0x06' function='0x0' />
</interface>
<input type='mouse' bus='ps2'>
    <alias name='input0' />
</input>
<input type='keyboard' bus='ps2'>
    <alias name='input1' />
</input>
<graphics type='vnc' port='5900' autoport='yes' listen='0.0.0.0' keymap='en-us'>
    <listen type='address' address='0.0.0.0' />
</graphics>
<video>

```



```
<model type='cirrus' vram='16384' heads='1' primary='yes'/>
<alias name='video0'/>
<address type='pci' domain='0x0000' bus='0x00' slot='0x02' function='0x0'/>
</video>
<memballoon model='virtio'>
  <alias name='balloon0'/>
  <address type='pci' domain='0x0000' bus='0x00' slot='0x04' function='0x0'/>
</memballoon>
</devices>
<seclabel type='dynamic' model='apparmor' relabel='yes'>
  <label>libvirt-faf9ab4e-57d0-4363-8ed2-15b5326918a5</label>
  <imagelabel>libvirt-faf9ab4e-57d0-4363-8ed2-15b5326918a5</imagelabel>
</seclabel>
<seclabel type='dynamic' model='dac' relabel='yes'>
  <label>+64055:+127</label>
  <imagelabel>+64055:+127</imagelabel>
</seclabel>
</domain>
```